

Diabetes Prediction using Logistic Regression

Import libraries

```
In [1]: import pandas as pd
import numpy as np
```

Import data frame

```
In [2]: df=pd.read_csv(r'https://github.com/YBI-Foundation/Dataset/raw/main/Diabetes.csv')
```

```
In [3]: df.head()
```

Out[3]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	age	diabetes
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

metadata:

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   pregnancies      768 non-null    int64
1   glucose          768 non-null    int64
2   diastolic        768 non-null    int64
3   triceps          768 non-null    int64
4   insulin          768 non-null    int64
5   bmi              768 non-null    float64
6   dpf              768 non-null    float64
7   age              768 non-null    int64
8   diabetes         768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [5]: `df.describe()`

Out[5]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.0
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.2
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.7
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.0
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.0
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.0
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.0
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.0

Getting columns:

In [7]: `df.columns`

Out[7]: Index(['pregnancies', 'glucose', 'diastolic', 'triceps', 'insulin', 'bmi', 'dpf', 'age', 'diabetes'], dtype='object')

In [10]: `df.shape`

Out[10]: (768, 9)

Get unique values(class/label) in y variables

In [8]: `df['diabetes'].value_counts()`

Out[8]: 0 500
1 268
Name: diabetes, dtype: int64

In [9]: `df.groupby('diabetes').mean()`

Out[9]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	age
diabetes								
0	3.298000	109.980000	68.184000	19.664000	68.792000	30.304200	0.429734	31.190000
1	4.865672	141.257463	70.824627	22.164179	100.335821	35.142537	0.550500	37.067100

Define y(dependent/label/target variable) and X (independent/feature/attribute variable)

```
In [50]: y=df['diabetes']
```

```
In [51]: y.shape
```

```
Out[51]: (768,)
```

```
In [52]: y
```

```
Out[52]: 0      1
         1      0
         2      1
         3      0
         4      1
         ..
        763     0
        764     0
        765     0
        766     1
        767     0
        Name: diabetes, Length: 768, dtype: int64
```

```
In [53]: X=df.drop(['diabetes'],axis=1)
```

```
In [54]: X.shape
```

```
Out[54]: (768, 8)
```

In [55]: X

Out[55]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 8 columns

Getting X variable standardized using MinMaxScaler

In [24]: `from sklearn.preprocessing import MinMaxScaler`In [25]: `mm=MinMaxScaler()`In [26]: `X=mm.fit_transform(X)`

In [27]: X

Out[27]: array([[0.35294118, 0.74371859, 0.59016393, ..., 0.50074516, 0.23441503,
0.48333333],
[0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,
0.16666667],
[0.47058824, 0.91959799, 0.52459016, ..., 0.34724292, 0.25362938,
0.18333333],
...,
[0.29411765, 0.6080402 , 0.59016393, ..., 0.390462 , 0.07130658,
0.15],
[0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,
0.43333333],
[0.05882353, 0.46733668, 0.57377049, ..., 0.45305514, 0.10119556,
0.03333333]])

Get model trained

Get probability of each predicted class

```
In [37]: lr.predict_proba(X_test)
```

```
Out[37]: array([[0.84199514, 0.15800486],
                [0.61583618, 0.38416382],
                [0.45354833, 0.54645167],
                ...,
                [0.74520246, 0.25479754],
                [0.85004994, 0.14995006],
                [0.24202189, 0.75797811]])
```

Get model evaluation

```
In [38]: from sklearn.metrics import confusion_matrix, classification_report
```

```
In [39]: print(confusion_matrix(y_test, y_pred))
```

```
[[333  13]
 [135  57]]
```

```
In [40]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.71	0.96	0.82	346
1	0.81	0.30	0.44	192
accuracy			0.72	538
macro avg	0.76	0.63	0.63	538
weighted avg	0.75	0.72	0.68	538

Get future predictions steps:

1. extract a random row using sample function
2. separate X and y
3. standardize X
4. predict

```
In [41]: X_new=df.sample(1)
```

In [42]: X_new

Out[42]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	age	diabetes
348	3	99	62	19	74	21.8	0.279	26	0

In [43]: X_new.shape

Out[43]: (1, 9)

In [44]: X_new=X_new.drop('diabetes',axis=1)

In [45]: X_new

Out[45]:

	pregnancies	glucose	diastolic	triceps	insulin	bmi	dpf	age
348	3	99	62	19	74	21.8	0.279	26

In [46]: X_new=mm.fit_transform(X_new)

In [47]: y_pred_new=lr.predict(X_new)

In [48]: y_pred_new

Out[48]: array([0], dtype=int64)

In [49]: lr.predict_proba(X_new)

Out[49]: array([[0.97981187, 0.02018813]])

predicted and actual class is zero(0) i.e. Non-Diabetic

In []: